

Watching Over the Reserve: A Fog-to-AWS Wildlife Pipeline

Hrishikesh Sajeev

Student ID: X25132377

MSc in Cloud Computing

Fog and Edge Computing (H9FECC)

National College of Ireland

x25132377@student.ncirl.ie

Abstract—Rangers patrolling a wildlife reserve on a fixed schedule learn about a poaching incident, a drying waterhole, or an unusual animal-movement spike only once they happen to pass that spot, not when it starts. Closing that gap is what the Java pipeline built for this paper does, from the sensors all the way to a live AWS account: two simulated reserves, five sensor types each, feed a plain-JDK HTTP fog node whose HTTP routes are discovered through Java reflection instead of a hand-written route table, and whose buffer is a lock-free, copy-on-write structure built on AtomicReference, not a synchronized block. On the backend, readings land first in an Amazon SQS queue instead of being processed inline; a Java 17 Lambda drains that queue and persists batched aggregates to DynamoDB, while a separate, independently deployed Lambda serves the dashboard's read endpoints over an Amazon API Gateway REST API by dispatching on a plain switch expression, without reflection or a third-party web framework. The complete system was deployed to an AWS account, where a DynamoDB Scan-pagination undercount, a missing SQS message-batching capability, and a credential-handling risk in a deploy script were found and fixed before the system went live. This paper documents that architecture, the alternatives weighed and set aside, and evidence gathered directly against the deployed system, not merely assumed from its source code alone.

Keywords—fog computing, edge computing, Internet of Things, serverless architecture, AWS Lambda, cloud deployment, wildlife conservation

I. INTRODUCTION

A ranger relying on a fixed patrol schedule finds out about a poaching incident, a waterhole running dry, or an unusual surge in animal movement only once the patrol happens to reach that part of the reserve, not when the event actually starts. Continuous sensing changes that trade-off: a reserve instrumented with always-on acoustic, motion, and environmental sensors can raise an alert the moment a threshold is crossed, without waiting for a human to walk past.

Fog computing extends that idea to where the data is produced. Bonomi et al. describe fog computing as a platform distinguished by low latency, wide geographic distribution, and support for a large number of edge nodes [1] -- properties that matter directly here, since a reserve's sensor network is exactly the kind of geographically spread, intermittently connected deployment the definition has in mind. The National Institute of Standards and Technology frames cloud computing around five essential characteristics: on-demand self-service, broad network access, resource pooling, rapid elasticity, and measured service [2]. Those characteristics describe the backend well but say nothing about what happens at the reserve itself, which is where a fog node earns its place in this design: it sits with the sensors, watches a rolling window of their readings, and is the one that decides -- via HabitatAlerts -- whether a threshold such as a spike in acoustic activity near a waterhole has been crossed. By the time anything reaches the cloud backend, an entire window of individual readings has already been reduced to a handful of summary figures plus whatever alerts fired, so the elastic, on-

demand side of the system in section II never has to ingest the raw sensor stream at all.

Three things had to come together for this paper's pipeline to count as done. The sensors and the fog node needed to actually buffer, aggregate, and evaluate five kinds of reading across two reserves at whatever sampling and dispatch rate each was told to run at, not a fixed rate baked into the code. The backend needed to turn that fog output into a dashboard capable of putting both reserves next to each other, using managed queueing and function-as-a-service rather than a server this project would have to keep patched and running itself. And none of that counted for anything unless it was deployed to a public cloud and checked against the deployment directly, not left as something the source code merely claimed to do.

The rest of the paper follows this structure. Section II lays out the architecture adopted at every layer, together with the alternatives that were weighed and set aside, and finishes with the security posture the deployed system settles into. Section III turns to implementation: the software stack and libraries in use, the automated test suite, the continuous-integration setup, and the AWS deployment itself, including the defects that surfaced once it went live. Section IV evaluates the deployed system using evidence measured from it. Section V closes with a critical assessment of the project and an outline of future work.

II. SYSTEM ARCHITECTURE AND DESIGN

A. Sensor and Fog Layer

Two simulated reserves, reserve-a and reserve-b, carry five sensor types apiece: camera-trap motion-event counts (motion_detection_count), an anomalous acoustic-signature level used as a proxy for gunshot or chainsaw detection (acoustic_poaching_risk_db), waterhole level (waterhole_level_cm), ambient temperature (ambient_temp_c), and soil moisture (soil_moisture_pct). Deploying low-cost acoustic sensors to flag gunshot-like signatures in a protected reserve is an established technique, not a hypothetical one: Prince et al. describe exactly this kind of algorithm running on open-source hardware in a tropical forest reserve [3], and acoustic_poaching_risk_db's threshold rule follows the same intent at a much smaller demonstration scale. Every sensor process reads two independently configurable environment variables, SAMPLE_INTERVAL and DISPATCH_INTERVAL, so sampling frequency and dispatch rate are decoupled knobs, not one value aliased to both.

Every reading is buffered in HabitatBuffer, whose entire state lives in a single AtomicReference<Map<FieldKey,List<Reading>>>, mutated exclusively through updateAndGet -- a lock-free, whole-structure copy-on-write retry loop, avoiding both a synchronized block and a per-key concurrent map. Each ingest() call builds a brand-new immutable snapshot map and hands it to updateAndGet(), which automatically retries the

supplied function if a competing thread's compare-and-swap won the race in between; `drainAll()` is a single `getAndSet(Map.of())` that atomically hands the whole retiring snapshot to the caller while resetting the live reference to empty. This compare-and-swap retry pattern for concurrent data structures is a well-studied technique, not an ad hoc choice: Michael and Scott's classic non-blocking queue algorithms establish the same underlying idea of retrying a compare-and-swap instead of blocking a competing thread [4]. The trade-off, acceptable at this demonstration's scale (five sensor types across two reserves, at most ten live keys), is that every `ingest()` copies the whole snapshot map; a 16-thread concurrent-ingest test proves no reading is ever lost to a missed retry.

On a fixed timer decoupled from buffer occupancy, `runWindowCycle()` calls `flushWindow()`, which drains `HabitatBuffer` through `drainAll()` and reduces each surviving key to five figures -- count, minimum, maximum, average, and the last-in-order latest reading -- before `HabitatAlerts.evaluate()` checks that figure set against Table I's threshold rules; `publishBatch()` then groups every window's aggregate-plus-alerts payload into `SendMessageBatch` calls capped at ten entries apiece, rather than one `sendMessage()` call per aggregate.

Sensor Type	Rule	Alert Key
acoustic_poaching_risk_db	avg > 75	poaching_risk_detected
waterhole_level_cm	avg < 20	drought_stress_risk
motion_detection_count	max > 30	unusual_activity_surge
soil_moisture_pct	avg < 10	habitat_dryness_risk

TABLE I. ALERT THRESHOLDS (EVALUATED ON THE WINDOW AGGREGATE)

HTTP routes on the fog node -- `/health`, `/thresholds`, and `/ingest` -- are discovered through Java reflection: every method annotated `@Route` on the gateway object is registered automatically by `AnnotatedRouter.bind()`, which scans the class's declared methods once at startup, with no hand-written route table to keep in sync. Reflection carries a documented runtime cost in exchange for that flexibility, a trade-off empirically characterised by Li et al. in their large-scale study of how Java applications actually use reflection [5]; at the request volumes this fog node handles, that cost is immaterial next to the benefit of never letting a route table drift out of sync with its handler methods.

B. Scalable Backend Layer

Nothing in the backend runs on a server this project provisions or patches itself -- every piece of it is a managed AWS primitive. `wcm-reserve-agg`, a standard Amazon SQS queue, sits between the fog node's dispatch rate and however fast the backend happens to be processing at a given moment. Messages arriving on that queue automatically invoke `wcm-processor`, an AWS Lambda function, via an event-source mapping, and each invocation writes one aggregated window into the `wcm-readings` DynamoDB table. Its partition key is `sensor_type`, but the sort key is not just `window_end` on its own -- it is the pair `window_end#site_id`, and the reason for pairing them only becomes clear once both reserves are reporting at the same time. If `reserve-a` and `reserve-b` each run a soil-moisture sensor and their flush cycles happen to close on the same `window_end`, a sort key built from `window_end` alone would give both aggregates an identical partition/sort combination, so the second `PutItem` call would simply replace the first with no error surfaced anywhere -- one reserve's

reading for that window would just be gone. Folding `site_id` into the sort key removes the ambiguity at the source, giving each reserve a separate row for the same window instead of requiring a global secondary index to tell them apart after the fact. The underlying pattern -- partition on a coarse attribute, disambiguate within it via the sort key -- is the same partition-and-replicate design DeCandia et al. describe as the basis for Amazon's internal Dynamo store, the system DynamoDB itself descends from [6].

Cold-start latency is the critique most often levelled at Lambda-based designs: without an already-warm execution environment, the runtime spends time initialising before the handler body even starts, and this delay has been measured directly in a large-scale production serverless deployment [7]. Two properties of this particular workload keep that latency from mattering in practice. The reading pipeline into `wcm-processor` is queue-triggered, so a slow cold start there simply delays how quickly a message is drained, not how quickly a browser renders anything -- there is no request waiting on it. The dashboard Lambda sidesteps the issue differently: constant polling from the live dashboard keeps its container active often enough that a true cold start is uncommon once a demonstration session has been running for a while.

C. A Switch-Based Dispatch Table for the Lambda Deployment

The dashboard-facing Lambda, `wcm-dashboard-api`, sits behind an Amazon API Gateway REST API and answers every `/api/*` route the dashboard calls: reserves, readings, thresholds, health, and backend-stats. Locally, the equivalent routes on `WildlifeDashboardApp` are discovered the same reflection-driven way the fog node's routes are, through `AnnotatedRouter` scanning `@Route`-annotated methods over a `com.sun.net.httpserver.HttpServer`. Both of those objects only exist for the lifetime of a running `HttpServer` instance, and a Lambda invocation never has one -- there is no listening socket, no `HttpExchange`, nothing for reflection to bind a handler to. Carrying `AnnotatedRouter` into the deployment as-is was considered and set aside for exactly that reason: doing so would have meant either an adapter translating API Gateway's event shape into a fake `HttpExchange`, or reimplementing the reflection scan against a different request type entirely, for a class of route table -- five fixed, static paths with no path parameters to extract -- that neither needs. `WildlifeDashboardLambda` instead answers each route through a plain Java switch expression on the concatenated method and path, calling directly into the same `ReserveRepository/PipelineChecks/ThresholdsGateway` classes the local `HttpExchange`-based routes already use, so the DynamoDB, SQS, and Lambda-metadata logic is never duplicated between the two entry points.

Four different things feed `wcm-dashboard-api` before it can answer a single request: readings and reserve comparisons come from the DynamoDB table, backlog depth comes from the SQS queue's own attributes, whether ingestion is still running comes from `wcm-processor`'s Lambda metadata, and fog-side status comes from the fog node's `/health` and `/thresholds` endpoints, reached over the Elastic IP its EC2 instance is bound to. That last one carries a practical caveat worth naming directly: reaching a bare IP address on port 8000 over unencrypted HTTP, rather than a named domain on 443, is precisely the pattern that corporate and campus firewalls are configured to flag and drop, so putting the dashboard itself behind that same IP would have made the whole page unreachable from exactly the kind of network a demo audience is likely sitting on. Its static files skip all of

that and sit in an Amazon S3 bucket instead, with dashboard.js resolving its own API base address at startup through an HTTP fetch of a small static/api-config.json resource generated fresh at each deploy -- not a build-time text substitution, a <meta> tag, or a separate config-global script -- left as an empty string for local development, where the same fog-node-adjacent host serves both origins.

D. Deployment Topology

The topology in Fig. 1 puts the two Lambda functions and everything downstream of them behind managed AWS services, leaving only the sensor and fog containers on infrastructure this project has to operate directly. Those

containers sit on a single EC2 instance with no key pair issued for it at all -- Session Manager is the only way in -- and a security group that blocks every port except the fog node's own 8000. From there the top half of the figure traces the ingest side left to right: sensors dispatch to the fog node, the fog node batches to SQS, SQS triggers wcm-processor, and wcm-processor writes into DynamoDB. The bottom half traces the query side: the browser loads the dashboard from S3 and, separately, calls API Gateway directly, which routes to wcm-dashboard-api. That second Lambda is the only component reading back across both halves, pulling from DynamoDB and the three other sources named above to answer every request the dashboard makes.




Fig. 1. Deployment topology: an ingest row (sensors through DynamoDB) and a query row (browser through the dashboard API), with the EC2-hosted sensor and fog tier the only piece not behind a managed AWS service.

E. Critical Analysis of Alternative Architectures

Several alternative designs were considered and either rejected or adopted for a specific, statable reason. HabitatBuffer's lock-free copy-on-write structure was itself a deliberate departure from a plain synchronized HashMap or a per-key ConcurrentHashMap: a whole-structure compare-and-swap retry is only correct here because the request volume at this scale never turns the snapshot-copy cost into a bottleneck, a trade-off that would not hold at a much larger request volume, where a per-key lock or a concurrent map would avoid copying the whole snapshot on every call.

A hand-written route table for the fog node's HTTP layer was considered in place of AnnotatedRouter's reflection scan, and rejected because it would have meant maintaining a second, parallel list of path-to-handler mappings that could silently drift out of sync with the actual @Route-annotated methods as endpoints were added during development -- a class of bug the reflection-driven approach makes structurally impossible, since the routing table is derived from the method metadata itself rather than copied from it by hand.

Apache Kafka was considered in place of SQS for the fog-to-processor link, and set aside. Kafka earns its complexity when several distinct consumer groups must independently rewind and re-read the same topic at their own pace, or when

partition-level ordering has to survive a consumer restart; wcm-processor is the sole reader of these messages, and DynamoDB already orders each item by its window_end sort key once written, so nothing downstream would ever call on a stream-replay guarantee. SQS needs no partitions to provision and no broker cluster to size, and its event-source mapping wires directly into Lambda without a consumer-group client to configure -- the simpler shape for a single, always-on consumer.

A single Lambda function handling both queue-triggered ingestion and API-Gateway-fronted query serving was considered in place of the two separate functions actually deployed, and rejected. wcm-processor's load tracks however many aggregate messages are waiting in the SQS queue at a given moment, while wcm-dashboard-api is invoked one-to-one by dashboard requests and must return within a much shorter window; merging the two under one reserved-concurrency limit and one timeout setting would mean a surge on either side eats into capacity the other needs -- a queue backlog could leave dashboard requests waiting behind a function busy draining messages, and a run of slow downstream reads on the query side could just as easily leave the queue undrained.

Finally, converting the fog node and sensors themselves to a scheduled-Lambda model, so that the entire system, not only

the backend, would be serverless, was considered and explicitly deferred rather than rejected outright. Lambda expects short, self-contained invocations with no state carried between them, which fits a queue consumer far better than it fits *HabitatBuffer*: the buffer's whole value comes from staying alive and accumulating readings between flushes, and turning that into a scheduled function would mean rebuilding the buffering logic around external storage instead of an in-process *AtomicReference*. This project's cloud-deployment scope was always the backend layer, not the sensor and fog layer, so that rebuild is left for Section V as future work rather than attempted here.

F. Security Considerations

Port 8000, reserved for the fog node, is the security group's sole inbound opening on this EC2 instance; nothing else gets in, including SSH, since the instance was never issued a key pair and all operator work instead happens over AWS Systems Manager Session Manager. That single open port exposes only read-only, non-sensitive status information, not data or control-plane access.

wcm-processor and wcm-dashboard-api both run under the Learner Lab account's one shared LabRole because this account will not let a student mint a new IAM role, not because a single broad role was ever the intended design. Split apart, wcm-dashboard-api would only ever need read access to DynamoDB, to the SQS queue's attributes, and to Lambda metadata, while wcm-processor's entire footprint would shrink to `sqs:ReceiveMessage`, `sqs:DeleteMessage`, and `dynamodb:PutItem` -- considerably tighter than what the shared role currently grants both functions by necessity. Leaving authentication off the dashboard's REST API was a deliberate choice, not an omission: every field it returns is an aggregated, non-personal reading from the reserve's sensors, so there is no sensitive payload for an authentication layer to be protecting in the first place.

III. IMPLEMENTATION

A. Software Components and Libraries

Java 17 runs across all four components -- sensor, fog, processor, and dashboard API -- with plain `com.sun.net.httpserver` standing in for a framework like Spring, and the official AWS SDK for Java v2 handling every AWS call. The overall pipeline shape -- sensors feeding a fog layer that windows, aggregates, and alerts before dispatching to a queue, a FaaS processor, a datastore, and a dashboard -- along with several of the implementation-choice axes below, carries over patterns already worked out in earlier project submissions, not built up from a blank page. None of that carried-over material comes from an earlier submission of this student's own; it originates in other students' codebases, and `readme.txt`'s REUSE section names each borrowed piece openly instead of leaving it unattributed. For every one of those axes -- the fog buffering mechanism, the alert-rule representation, the SQS publisher shape, the HTTP routing style, and the sensor-loop scheduling mechanism -- the actual present-day source it was adapted from was opened and read line by line while writing this report, and each was reshaped into something that departs from that source in a concrete, identifiable way, not copied from it outright. The domain-specific code (sensor types, thresholds, the field-log logic, and the entire dashboard UI) is original to this project.

Jackson (`jackson-databind`) handles JSON throughout: parsing `/ingest` payloads, rendering `/thresholds`, and building the fog-to-SQS aggregate-plus-alerts message body. `Chart.js`

draws the dashboard's waterhole-level trend chart client-side, vendored locally instead of pulled from a content-delivery network, so the page carries no third-party runtime dependency. Local development and continuous integration both run Docker Compose against LocalStack, an open-source AWS cloud emulator, so the same SQS, DynamoDB, and Lambda interactions exercised in production can be run day to day without touching AWS itself.

B. Testing Strategy

Each Maven module carries its own JUnit 5 test suite: 7 tests for sensors, 44 for the fog layer, 5 for the backend processor, and 26 for the dashboard, all currently passing. `HabitatGatewayHttpTest` and `AnnotatedRouterTest` exercise `/ingest` and the reflection-driven route dispatch over a real `com.sun.net.httpserver.HttpServer` bound to an ephemeral port, not a unit test of validation logic in isolation. `HabitatBufferTest` includes a 16-thread concurrent-ingest test proving the compare-and-swap retry loop never drops a reading under real contention. `ThresholdsGatewayTest` covers both the success path and an unreachable-upstream path for the dashboard's fog-thresholds proxy.

Two tests specifically target the defects described in the next subsection. `PipelineChecksTest` and `WildlifeDashboardLambdaTest` each assert that a four-page fake DynamoDB scan (400, 400, 400, and 87 items) sums to exactly 1287 through `itemCount()`'s `scanPaginator`-based pagination, not just the first page's count. `ReservePublisherTest` asserts that a 23-message flush window batches into `SendMessageBatch` calls of size ten, ten, and three, not twenty-three individual `sendMessage()` calls.

C. Continuous Integration and Deployment

Every push and pull request triggers a two-stage GitHub Actions pipeline, not a single monolithic job. Stage one, unit-tests, executes the Maven suite across all four Java modules in complete isolation -- no containers spun up, no socket ever opened. Stage two, integration, is gated on stage one succeeding and only then launches the LocalStack-backed `docker-compose.yml` environment, drives `infra/verify_pipeline.py` through it end to end, and tears the whole stack back down once finished, pass or fail, so no container is ever left running for the next job to trip over.

D. AWS Deployment and Defects Discovered

Nothing in this stack stayed a local simulation by the end: an AWS Academy Learner Lab account in us-east-1 now hosts every piece of it, sensor and fog containers included, alongside the queue, both Lambda functions, DynamoDB, API Gateway, and the S3 frontend, and that shift from emulator to a live account is exactly what turned up three problems Docker Compose and LocalStack alone never would have.

First, the number reported by wcm-dashboard-api's backend-stats endpoint as the row count in DynamoDB came from a single `Select.COUNT` call to `Scan`. What that call actually returns is not the table's true size but whatever fits in roughly 1 MB of scanned data; DynamoDB signals there is more by attaching a `LastEvaluatedKey` to the response, and unless the caller issues another `Scan` starting from that key, the number reported stops short of reality with nothing to flag the shortfall. The bug is baked into how `Scan` itself works, not something that only surfaces at a particular scale, and it was caught and corrected here before the table ever grew large enough to expose it visibly. The fix uses the AWS SDK's `scanPaginator()`, an `Iterable` that follows `LastEvaluatedKey`

across pages automatically, rather than a hand-rolled while loop, a do-while loop, or a recursive call.

Second, the fog node's flush cycle originally sent one SQS message per closed (sensor_type, site_id) group in a loop -- functionally correct, but wasteful whenever more than one group closed in the same flush window, the normal case for a two-reserve, five-sensor-type deployment. The fix, `ReservePublisher.publishBatch()`, collects every window's aggregate into one list and chunks it at `SendMessageBatch`'s ten-entry limit, and `HabitatGateway.runWindowCycle()` now calls it once per flush cycle instead of looping the earlier single-message `publish()` call.

Third, `backend/processor/deploy_lambda.sh -- tooling` written specifically for deploying against LocalStack -- defaults to the LocalStack endpoint whenever `AWS_ENDPOINT_URL` is unset, and unconditionally hardcodes a test/test static credential pair regardless of that variable's value. Left unguarded, this would have meant the script silently overwrote any real AWS credentials in the environment if it were ever pointed at a live account by mistake. This was not exercised in the actual deployment -- the Lambda functions were created directly through the AWS CLI with their own explicit environment variables, entirely bypassing this script -- but the risk was documented directly in the script with a comment clarifying it is LocalStack-only tooling, so a future reader does not assume it is safe to run against a live account.

None of the three fixes were taken on faith from the unit test suite; each was checked again against the running deployment. For the pagination and batching fixes, that meant the specific assertions already described -- the four-page sum and the ten/ten/three chunk split -- passing against the fakes stand in for what the actual DynamoDB and SQS clients do. For the dashboard, it meant loading the deployed page in an actual browser and reading its network requests directly: every static asset resolved, and the live data climbed across repeated reloads instead of sitting on a cached, stale count. Setup steps and the deployment history summarised above are recorded in the project's `readme.txt`; the source itself, fixes included, is hosted at [GitHub repository URL].

IV. EVALUATION

A. Automated Test Coverage

Table II summarises the automated test suite as of the final verified deployment. All 82 tests pass against both the LocalStack-backed development environment and, for the parts of the suite that exercise genuine client-construction logic, against the fixes confirmed on the AWS deployment described in Section III-D.

Module	Tests	Focus
sensors	7	bounded random walk, dual-rate config
fog (buffer+aggregate+alerts)	18	lock-free buffer, windowed aggregation, threshold rules
fog (publisher, routing, HTTP)	26	SQS batching, reflection routing, real-socket /ingest
backend/processor	5	DynamoDB item shape, sort-key logic
backend/dashboard	26	routes, pagination fix, real-socket + Lambda dispatch tests

TABLE II. AUTOMATED TEST SUITE BY MODULE

B. Live Deployment Health

Independent verification did not stop at the deployed API's `/api/health` response, though that response was checked too and found consistent: gateway (the fog node's health check), queue, lambda, and pipeline every field true, `freshness_age_seconds` under six seconds. Each underlying AWS resource was then queried directly with the AWS CLI, separately from anything the application itself reported: a get-function call on both Lambdas returned `State Active` with `LastUpdateStatus Successful`, get-event-source-mapping for the queue feeding `wcm-processor` came back `State Enabled`, `describe-instances` showed the EC2 host running and bound to its Elastic IP with a security group locked to inbound TCP 8000 only and no key pair attached, and `describe-table` on DynamoDB returned `ACTIVE` under on-demand billing.

Checking the `backend-stats` item count twice, fifteen seconds apart, showed it climb from 362 to 378 -- evidence the pipeline was actively running, not serving a frozen or cached figure. The dashboard itself was opened in a real browser, not just probed with `curl`: the page, its stylesheet, its vendored chart library, and `dashboard.js` each came back with a 200 and no console errors, and both `reserves`' field-station log panels updated with live, changing data, correctly tripping alert flags against the threshold rules in Table I.

C. Cost Considerations

Every backend component other than the EC2 instance runs within AWS's request-based free-tier allowances at this project's data volume: DynamoDB on-demand billing charges per request rather than per provisioned throughput, both Lambda functions' invocation counts sit well inside the perpetual one-million-request Lambda free tier, and API Gateway and SQS are billed the same way, per request, with no idle cost between requests. The EC2 instance is the only continuously billed resource, since it runs the fog node and sensor containers as long-lived processes, not on-demand invocations. DynamoDB, Lambda, and API Gateway are what actually keep the dashboard and its API answering requests, so powering the EC2 instance off between demo sessions does not break either one -- what stops working is narrower than that: the health check's fog-reachability flag goes false, and no new readings land in the table until the instance comes back up.

D. Limitations

This evaluation has four honest gaps, named here rather than left for a reader to find unaided. The fog and sensor tier is the first: it still runs on one EC2 instance, so that specific tier keeps a single point of failure even though everything downstream of it is redundant AWS infrastructure; Section II-D's critical analysis explains why moving this tier to a fully serverless model was deferred, not attempted, but that does not make the resulting availability gap any less real. Second, this project ran inside an AWS Academy Learner Lab account, which is provisioned for bounded class sessions, not the kind of uninterrupted, multi-day run a production rollout would need to prove itself over, so every figure in this section reflects one continuous sitting, not sustained operation across days or weeks. Third, one broadly scoped IAM role is shared between both Lambda functions instead of each holding a least-privilege role, a constraint imposed by the Learner Lab account, discussed further in Section II-F, and not a deliberate security choice. Fourth, the automated integration suite verifies the pipeline against LocalStack, not the live AWS account itself, so the live evidence in this section rests on

manual polling and browser inspection, not on an automated live-account test run.

V. CONCLUSION

Wildlife habitat monitoring was the domain chosen to prove out a fog-to-serverless pipeline end to end: sensors and a fog layer with genuine configurability, a backend split across two Lambda functions rather than one, and an actual cloud deployment to stand behind the design choices made at every layer, not just a description of them on paper. By the close of this work all three layers existed as running, tested, deployed artefacts, each weighed against the alternatives that were realistically on the table, rather than a paper design left untested.

Two defects made it through the full local test suite and a complete LocalStack-backed integration run without tripping a single assertion: the DynamoDB pagination undercount, and the missing SQS batching. Neither one changes what a single-page scan or a single-message send returns at this project's small local data volume -- they only change what gets silently dropped once a real account's table or a flush cycle's group count grows past that boundary, and no local test exercises data at that scale. What actually caught both was pointing the same code at the deployed AWS account and at the rendered dashboard in a real browser, not writing more local tests.

Application code was not where the actual difficulty sat. It was in the AWS Academy Learner Lab account's own platform and IAM restrictions, and in resisting the temptation to treat a green API check as evidence that the page a user actually sees was working too. Two extensions would matter most going forward: moving the fog and sensor tier off a continuously running EC2 instance onto a fully event-driven, scheduled-Lambda model, and, were the account's IAM restrictions ever lifted, splitting the shared LabRole into two least-privilege roles -- one for wcm-processor, one for wcm-dashboard-api.

VI. REFERENCES

- [1] F. Bonomi, R. Milito, J. Zhu, and S. Addepalli, "Fog computing and its role in the Internet of Things," in Proc. 1st ACM Mobile Cloud Computing Workshop (MCC '12), Helsinki, Finland, Aug. 2012, pp. 13–16.
- [2] P. Mell and T. Grance, "The NIST Definition of Cloud Computing," National Institute of Standards and Technology, Gaithersburg, MD, USA, Special Publication 800-145, Sep. 2011.
- [3] P. Prince, A. Hill, E. Piña Covarrubias, P. Doncaster, J. L. Snaddon, and A. Rogers, "Deploying acoustic detection algorithms on low-cost, open-source acoustic sensors for environmental monitoring," *Sensors*, vol. 19, no. 3, p. 553, 2019.
- [4] M. M. Michael and M. L. Scott, "Simple, fast, and practical non-blocking and blocking concurrent queue algorithms," in Proc. 15th Annu. ACM Symp. Principles of Distributed Computing (PODC '96), Philadelphia, PA, USA, May 1996, pp. 267–275.
- [5] Y. Li, T. Tan, and J. Xue, "Understanding and analyzing Java reflection," *ACM Trans. Softw. Eng. Methodol.*, vol. 28, no. 2, Art. 7, Feb. 2019.
- [6] G. DeCandia et al., "Dynamo: Amazon's highly available key-value store," in Proc. 21st ACM SIGOPS Symp. Operating Systems Principles (SOSP '07), Stevenson, WA, USA, Oct. 2007, pp. 205–220.
- [7] M. Shahrad et al., "Serverless in the wild: Characterizing and optimizing the serverless workload at a large cloud provider," in Proc. 2020 USENIX Annu. Tech. Conf. (USENIX ATC '20), Jul. 2020, pp. 205–218.