

# A Fog Computing Pipeline for Data Center Environmental Monitoring: Design, Deployment, and Evaluation on AWS

Nithin

Student ID: X25125338

MSc in Cloud Computing

Fog and Edge Computing (H9FECC)

National College of Ireland

x25125338@student.ncirl.ie

**Abstract**—This paper presents the design, implementation, and real-world cloud deployment of a fog and edge computing pipeline for data center environmental monitoring. The system simulates five distinct sensor types (temperature, humidity, airflow, power load, and dust density) across two server halls, each independently configurable for sampling and dispatch frequency. A Node.js fog node buffers incoming readings in a fixed-size ring buffer, performs windowed aggregation and threshold-based alerting, and dispatches each window's results as a batched message. The backend is entirely serverless and split across two independently deployed AWS Lambda functions. The first is triggered by an Amazon SQS queue and writes aggregated windows to Amazon DynamoDB, while the second performs internal request routing and sits behind a real Amazon API Gateway REST API. This design avoids a dashboard server that calls AWS services directly. The dashboard's static assets are served from Amazon S3 and call the API Gateway endpoint, with no intermediary server. The complete system was deployed and tested on a real AWS account instead of relying solely on a local emulator, which surfaced defects invisible to local testing. These included credential-handling bugs that only manifest under temporary AWS credentials and a DynamoDB pagination bug that silently undercounted large table scans. This paper critically analyses the architectural decisions taken at each layer, the defects discovered during real deployment, and the lessons learned for building cloud-native fog computing systems.

**Keywords**—fog computing, edge computing, Internet of Things, serverless architecture, AWS Lambda, cloud deployment, data center monitoring

## I. INTRODUCTION

Data centers depend on continuous environmental monitoring to protect the hardware they house. An undetected overheating server, a humidity swing that risks condensation on sensitive equipment, restricted airflow around a rack, an unexpected power-load spike, or rising dust density between scheduled inspections can each trigger a hardware failure or a costly unplanned outage. This is precisely the class of condition wireless sensor deployments in real data centers have been built to catch before it escalates [1], [2]. A facility relying on a fixed inspection schedule instead of continuous sensing discovers these conditions only after the fact, precisely when a condition-driven system is meant to surface them immediately. Internet of Things sensing, combined with fog and edge processing close to the sensors themselves, allows a facilities team to evaluate every reading against explicit threshold rules as it arrives; there is no need to wait for the next scheduled walk-through.

The National Institute of Standards and Technology defines cloud computing around five essential characteristics: on-demand self-service, broad network access, resource

pooling, rapid elasticity, and measured service [3]; fog computing extends that model instead of replacing it. Bonomi et al. define fog computing as a platform that extends cloud services to the edge of the network, characterised by low latency, geographic distribution, and support for a large number of nodes [4]. This project adopts that definition: windowed aggregation and threshold evaluation are performed on a fog node co-located with the sensors, and only the resulting aggregate-plus-alerts payload, not every raw reading, crosses the network boundary into the cloud.

This paper documents the design and construction of a fog and edge computing pipeline for this domain. The system was designed to satisfy three requirements at minimum. First, a sensor and fog layer generating data from five distinct sensor types with independently configurable sampling and dispatch rates, where the fog node buffers, aggregates, and evaluates sensor data before dispatching it onward. Second, a scalable backend layer, built from managed queue and function-as-a-service mechanisms, that processes the fog node's output and exposes a responsive dashboard for the sensor types. Third, deployment and testing of that backend on a public cloud platform, not merely a local simulation.

The remainder of this paper is organised as follows. Section II describes the system's architecture and the design decisions taken at each layer. This includes the separately deployed API backend that distinguishes this project's architecture, a critical comparison against the alternatives that were not chosen, and the security posture of the resulting deployment. Section III describes the concrete implementation: the software components and libraries used, the automated test suite, the continuous integration pipeline, and the real AWS deployment together with the defects it surfaced. Section IV evaluates the deployed system against measured evidence rather than intent alone. Section V concludes with a critical reflection on the project and directions for future work.

## II. SYSTEM ARCHITECTURE AND DESIGN

### A. Sensor and Fog Layer

The simulated deployment models two server halls, hall-1 and hall-2, each instrumented with five sensor types: temperature\_c, humidity\_pct, airflow\_cfm, power\_load\_kw, and dust\_density\_ugm3. Every one of the ten resulting sensor processes exposes two independent, separately configurable environment variables, SAMPLE\_INTERVAL and DISPATCH\_INTERVAL, so that generation frequency and dispatch rate are decoupled knobs rather than an aliased value.

The fog node buffers incoming readings in a fixed-size ring buffer keyed by (sensor\_type, site\_id). Each key holds a

plain array of RING\_CAPACITY (256) slots with a manually tracked write index that advances modulo the capacity, a bounded circular-buffer structure for streaming data in place of an unbounded growing list [5], [6]. On every window tick, the fog node reconstructs each ring's contents in oldest-first order and computes the window's minimum, maximum, average, latest value, and reading count. It then evaluates the aggregate against the domain-specific threshold rules listed in Table I and resets the ring for the next window. The resulting aggregate-plus-alerts payload is dispatched to an Amazon SQS queue through an EventEmitter-decoupled publisher. Window-close and the SQS send remain two separate concerns joined only by an emitted event, echoing the general publish-subscribe pattern of separating event producers from consumers in space, time, and synchronisation [7], [8]. The publisher always issues a batched SendMessageBatchCommand, chunked at SQS's ten-entry ceiling, so every group closed in one flush cycle goes out together whenever more than one is ready.

TABLE I. ENVIRONMENTAL ALERT RULES

| Sensor Type       | Rule           | Alert Key             |
|-------------------|----------------|-----------------------|
| temperature_c     | avg > 27°C     | overheat_risk         |
| humidity_pct      | avg > 60%      | condensation_risk     |
| humidity_pct      | avg < 20%      | static_discharge_risk |
| airflow_cfm       | avg < 400 CFM  | insufficient_cooling  |
| power_load_kw     | avg > 130 kW   | capacity_warning      |
| dust_density_ugm3 | avg > 50 ug/m3 | air_quality_risk      |

### B. Scalable Backend Layer

The backend layer is built entirely from managed, scalable AWS primitives. An Amazon SQS standard queue (dce-hall-agg) decouples the fog node's dispatch rate from the backend's processing rate. An AWS Lambda function (dce-processor), subscribed to that queue via an event-source mapping, is invoked automatically as messages arrive, writing each aggregated window into an Amazon DynamoDB table (dce-readings). The table uses sensor\_type as its partition key and a composite sort key of the form window\_end#site\_id as its range key. This follows the same partition-key philosophy behind Amazon's own eventually-consistent, highly available key-value store: a narrow, well-known set of key-based lookups rather than ad-hoc relational queries [9]. This design choice is deliberate: if the sort key were window\_end alone, hall-1 and hall-2 producing an aggregate for the same sensor type in the same flush cycle would collide on an identical primary key. Whichever write landed second would silently overwrite the first. Appending the site identifier to the sort key disambiguates the two halls' concurrent writes without requiring any additional index.

A commonly raised critique of Lambda-based architectures is cold-start latency, an initialisation delay incurred whenever an invocation finds no currently warm execution environment, a cost documented empirically across commercial FaaS platforms [10]. This is a cost of the FaaS approach, but it is acceptable here for the same reason it is acceptable for the ingestion path: the fog-to-processor path is asynchronous and queue-triggered, not a user-facing request/response path. An occasional cold-start delay when a reading is durably stored is therefore not a user-visible page load.

### C. A Separately Deployed API Backend Behind a Real API Gateway

Unlike a design in which the dashboard's backend process calls DynamoDB, SQS, and Lambda directly, this project's individually required architecture deploys a second, entirely separate Lambda function, dce-api. The function performs internal path-and-method routing against an ordered [method, regex, handler] table, which can be unit-tested with fake AWS clients and no Lambda runtime at all. dce-api is fronted by a real Amazon API Gateway REST API: a {proxy+} resource under the root, with an ANY-method AWS\_PROXY integration. Every request the dashboard makes to any /api/\* path is therefore proxied verbatim to dce-api, which then decides internally how to answer it. This centralises cross-cutting request handling behind a single gateway rather than exposing each backend capability as a separately reachable endpoint [11]. dce-api answers five endpoints: GET /api/readings, GET /api/halls (and /api/halls/:hallId for an individual hall), GET /api/health, GET /api/backend-stats, and GET /api/thresholds. The last of these is proxied from the fog node's /thresholds endpoint, letting a client retrieve the live alert rules without duplicating them.

### D. Dashboard Layer and Its Real-AWS Adaptation

The dashboard was designed from the outset as a deliberately minimal reverse proxy, not a REST API. Its Node.js HTTP server resolves the API Gateway's invoke URL exactly once at startup and forwards every /api/\* request to it verbatim, holding no AWS SDK imports and no route table. Locally, this same server also serves the dashboard's static assets on one port. For the real AWS deployment, this reverse-proxy role was retired rather than ported to a second EC2 process. The dashboard's HTML, CSS, and JavaScript are now served from an Amazon S3 bucket over the bucket's HTTPS endpoint, and the frontend calls the API Gateway's invoke URL directly. This URL is configured through a meta tag in place of a hard-coded value, and the dce-api Lambda's JSON responses carry an Access-Control-Allow-Origin header for the resulting cross-origin request.

This mirrors a constraint discovered during a separate project in this portfolio: a raw EC2 public IP address on a non-standard port served over plain HTTP is exactly the traffic shape that campus and corporate WiFi networks commonly block by policy. Hosting the dashboard on EC2 at all would therefore have reintroduced a reachability risk that the fully serverless path avoids entirely.

### E. Deployment Topology

Fig. 1 summarises the resulting topology. Sensor processes and the fog node run in Docker containers on an EC2 instance, managed entirely through AWS Systems Manager Session Manager in place of SSH. As a result, no key pair is provisioned, and no inbound SSH port is opened on the instance's security group. The fog node dispatches aggregated windows to an SQS queue, which triggers dce-processor, which writes to DynamoDB. dce-api, sitting behind API Gateway, reads from that same table (and from the queue, dce-processor's metadata, and the fog node's health/thresholds endpoints) to answer the dashboard's REST API calls. The dashboard's static frontend, by contrast, is served independently from Amazon S3 and reached directly by the browser. This topology places every backend component except the sensor and fog tier behind a managed, serverless service boundary.




Fig. 1. Deployment topology: sensors and the fog node run on EC2; the backend (SQS, both Lambda functions, DynamoDB, API Gateway) is fully serverless; the frontend is served from S3.

#### F. Critical Analysis of Alternative Architectures

Several alternative designs were considered and rejected, each for a specific, statable reason. A single Lambda function handling both queue-triggered ingestion and API-Gateway-fronted request serving was considered in place of two separate functions, and rejected. Ingestion and query serving have different invocation triggers and different concurrency and timeout characteristics. Coupling them in one function would mean a slow or failing query path could starve the ingestion path of concurrency, or vice versa. This is exactly the failure mode that splitting the two functions is meant to avoid.

An unbounded, per-key growing array was considered in place of the fixed-size ring buffer for fog-side buffering, and rejected in favour of a bounded structure: an unbounded buffer risks unconstrained memory growth if a downstream flush is ever delayed. A fixed-capacity ring buffer with a wrapping write index guarantees a hard upper bound on memory use regardless of how long a window runs, at the cost of silently discarding the oldest readings once that bound is exceeded. This is an accepted trade for a workload whose window length is short relative to its sampling rate.

A synchronous dashboard backend that calls AWS services in-process, the pattern the rest of this portfolio's Node.js siblings use, was considered in place of the decoupled, API-Gateway-fronted dce-api Lambda. That approach was rejected specifically because this project's individual rubric

required an architecture that differs structurally, not merely superficially, from an always-running backend process that imports the AWS SDK directly.

Finally, converting the fog node and sensors themselves to a scheduled-Lambda model, so that the entire system, not only the backend, would be serverless, was considered and explicitly deferred rather than rejected outright. A continuous sensor loop and a stateful ring buffer both map awkwardly onto Lambda's stateless, time-limited invocation model without a redesign of the buffering logic itself. This project's cloud-deployment scope is also limited to the backend layer, not the sensor and fog layer. Pursuing this conversion was judged a larger redesign than the remaining project time justified, and is recorded as future work in Section V.

#### G. Security Considerations

The EC2 instance is administered exclusively through AWS Systems Manager Session Manager; no key pair was ever provisioned, and its security group has no inbound rule for port 22, so there is no SSH attack surface at all. The only inbound rule opened is for the fog node's health and threshold endpoints (port 8000), which dce-api needs to reach, and it exposes only read-only, non-sensitive status information, not data or control-plane access.

Both Lambda functions execute under the AWS Academy Learner Lab account's shared LabRole, since this account's service-control policies prevent a student from creating new, more narrowly scoped IAM roles. This is an actual limitation,

not a design preference: in a production account, dce-processor and dce-api would each hold their own least-privilege role (dce-processor needing only `sqs:ReceiveMessage/DeleteMessage` and `dynamodb:PutItem`; dce-api needing only read-oriented `DynamoDB`, `SQS`, and `Lambda-metadata` actions). This gap is recorded honestly here, not glossed over. The dashboard's REST API is intentionally unauthenticated, since it exposes only aggregated, non-personal environmental readings for a facility, not data for which access control would be a meaningful requirement at this project's scale.

### III. IMPLEMENTATION

#### A. Software Components and Libraries

The sensor, fog, processor, and dce-api components are all implemented in Node.js 20 using plain CommonJS (no TypeScript build step) and the official AWS SDK for JavaScript v3 for every AWS interaction. In place of a web framework such as Express, they rely on Node's built-in http module, keeping the dependency surface deliberately small. The dashboard's trend visualisation is rendered client-side with Chart.js. The library is vendored locally, not loaded from a content-delivery network, so the page carries no third-party runtime dependency. Local development and continuous integration both use Docker Compose together with LocalStack, an open-source AWS cloud emulator, so the same SQS, DynamoDB, and Lambda interactions used in production can be exercised without touching a real AWS account during day-to-day development.

#### B. Testing Strategy

Every module has an automated test suite written with Node's built-in test runner, exercised through hand-written fake AWS client objects rather than mocking libraries. Each unit test therefore asserts on the exact request shape sent to a given AWS SDK call. The ring buffer's test suite specifically covers wraparound behaviour, including a case where writes exceed capacity mid-window and only the last `RING_CAPACITY` readings survive. It also verifies that every ring's slots, write index, and count are reset to empty on every window tick. The EventEmitter-decoupled publisher is tested for exactly one `SendMessageBatchCommand` call carrying multiple Entries when more than one group closes in the same flush cycle. dce-api's internal router is tested by dispatching method-and-path-regex matches against fake AWS clients directly, with no Lambda runtime or API Gateway involved at all.

Beyond unit tests, a load-testing script bursts a configurable number of synthetic messages onto the real queue. It makes two assertions: first, that the queue depth is non-zero immediately after the burst, confirming the messages reached SQS and were not silently dropped; second, that the queue depth has either fully drained or measurably decreased within a timeout window. This second, timeout-tolerant check allows for the underlying Lambda concurrency ceiling without treating a slower-than-expected drain as a hard failure.

#### C. Continuous Integration and Deployment

A GitHub Actions workflow runs the full automated test suite for all five modules on every push to the repository. A second job then brings the entire stack up with Docker Compose against LocalStack and runs a dedicated end-to-end verification script that exercises the live pipeline, submits synthetic readings, and asserts that they land in DynamoDB with the expected `sort_key` shape. The stack is torn down

afterwards regardless of outcome, so a failed run cannot leave a stale container running for the next one.

A second, separately triggered GitHub Actions workflow exists specifically to restart the EC2-hosted fog and sensor tier once an AWS Academy Learner Lab session ends and a fresh one begins, since the Learner Lab platform stops the EC2 instance between sessions. It reads its AWS credentials from encrypted GitHub repository secrets and verifies that the credentials resolve to the expected AWS account before doing anything else. Only then does it start the instance, wait for its Systems Manager agent to register, and re-verify that the fog node is reachable both directly and through dce-api's gateway field before reporting success.

#### D. Real AWS Deployment and Defects Discovered

The complete system, including the sensor and fog tier, the SQS queue, both Lambda functions, the DynamoDB table, the API Gateway endpoint, and the S3-hosted frontend, was deployed to a real AWS account (an AWS Academy Learner Lab account) in region `us-east-1`. As a result, the backend was deployed and tested on an actual public cloud platform rather than only a local emulator. This real deployment surfaced two classes of defect that the local, LocalStack-backed test suite had not caught.

First, three separate files across the fog, processor, and dce-api components constructed their AWS SDK client credentials using a check that was true in both the local and the real-AWS case. Whenever an AWS access key environment variable was present, the code applied a hardcoded, LocalStack-only static credential pair; one file did so unconditionally, with no check at all. Because real AWS Lambda and EC2 execution environments always populate that same environment variable automatically, with temporary, session-token-bearing credentials, this logic silently discarded the real temporary credentials in favour of the hardcoded pair on every real deployment. This would have caused every DynamoDB and SQS call to fail authentication in production, despite passing every LocalStack-backed test. The fix was to gate the hardcoded override on the presence of an explicit LocalStack endpoint override instead, the only condition actually unique to the local-emulator case. Each of the three files expressed this gating logic with a different code shape, not the same patch copy-pasted three times.

Second, dce-api's backend-stats endpoint reported the total number of items stored in DynamoDB via a `COUNT-select` Scan call. A `COUNT-select` scan only counts the single page it reads, typically capped around 1MB of table data; a table larger than that returns a `LastEvaluatedKey` that must be followed with a further scan, or the reported count silently undercounts the table. This defect was found and fixed before it manifested as a visible symptom in this project's data volume. The same pagination gap had already caused a real, deployment-triggered undercount in a separate project in this portfolio built against the same AWS SDK. The fix follows `LastEvaluatedKey` across pages and sums every page's `Count`, with two new tests covering the multi-page-summing and scan-failure cases.

Both defects were fixed and re-verified against the live deployment: the credential fix was confirmed via a DynamoDB write and SQS receive succeeding end to end after redeployment. The pagination fix was confirmed by a real `items_in_table` count returned correctly against the deployed table, rather than trusted from the unit test suite alone. The source code, including the fixes described above, is available at <https://github.com/nithinveshala/fog-and-edge>,

together with a readme.txt documenting installation, configuration, and the deployment history described in this section.

#### IV. EVALUATION

##### A. Automated Test Coverage

Table II summarises the automated test suite as of the final verified deployment. All 114 tests pass against both the LocalStack-backed development environment and, for the parts of the suite that exercise real client-construction logic, against the fixes verified against the real AWS deployment described in Section III-D.

TABLE II. AUTOMATED TEST SUITE BY MODULE

| Module            | Tests | Focus                                                   |
|-------------------|-------|---------------------------------------------------------|
| sensors           | 12    | dual-rate sampling/dispatch independence                |
| fog               | 46    | ring buffer wraparound, windowed aggregation, alerting  |
| backend/processor | 9     | DynamoDB item shape, sort-key logic                     |
| backend/api       | 37    | internal routing, DynamoDB pagination, thresholds proxy |
| backend/dashboard | 10    | reverse-proxy behaviour, invoke-URL resolution          |

##### B. Live Load Test

A burst load test sent 300 synthetic messages directly to the real, deployed dce-hall-aggr queue in 5.47 seconds (approximately 55 messages per second). The test used synthetic sensor-type names distinct from the five real sensor types, so the burst could never land in the DynamoDB partitions the live dashboard reads from. Amazon SQS's `ApproximateNumberOfMessages` and `ApproximateNumberOfMessagesNotVisible` attributes reported a queue depth of zero immediately after the burst completed, a result that would ordinarily read as a failed send. A direct DynamoDB query against the five synthetic sensor-type partitions, however, confirmed that all 300 messages had been received, processed, and written by dce-processor before the follow-up check even ran.

This is consistent with a documented property of real SQS rather than evidence of a broken pipeline. The queue's own approximate-count attributes are explicitly eventually consistent and can under-report immediately after a burst, a behaviour that LocalStack's more immediately consistent emulation does not reproduce. This evaluation would not have discovered that behaviour without testing against the real service.

##### C. Live Deployment Health

Polling the deployed API's `/api/health` endpoint twice, fifteen seconds apart, returned `freshness_age_seconds` of 7.429 and 3.357 respectively, and `items_in_table` counts consistent with a live, continuously updating pipeline rather than static or cached data. All AWS resources (the DynamoDB table, the SQS queue, both Lambda functions, the API Gateway API, and the EC2 instance) were independently confirmed healthy via direct AWS CLI calls, not merely trusted from the application's self-reported status.

Operationally, every backend component other than the EC2 instance runs within AWS's request-based free-tier allowances at this project's data volume. The EC2 instance itself is the only continuously billed resource, and it can be stopped between demonstration sessions without affecting the dashboard or its API. Neither the dashboard nor its API

depends on the instance being available, except for the gateway field of `/api/health` and freshly generated sensor data.

##### D. Limitations

Three limitations of this evaluation are worth stating explicitly instead of leaving them implicit. First, the fog and sensor tier still runs on a single EC2 instance and therefore remains a single point of failure for that tier specifically, even though the backend behind it is fully redundant AWS-managed infrastructure. The critical analysis in Section II-F explains why converting this tier to a fully serverless model was deferred, not attempted, but the resulting availability gap is real and unresolved, not merely theoretical. Second, the AWS Academy Learner Lab account used for this deployment is session-based, not persistent, so the sustained, multi-day uptime evaluation a production deployment would warrant was not possible within the account's constraints. The evaluation in this section reflects verified behaviour within one continuous session, not long-run reliability. Third, both Lambda functions share one broadly scoped IAM role instead of two least-privilege roles, a constraint of the Learner Lab account and not a deliberate security design, as discussed in Section II-G. This is an actual limitation of what could be independently verified about the deployment's security posture, not evidence that the least-privilege alternative was evaluated and rejected.

#### V. CONCLUSION

This project set out to build a complete fog and edge computing pipeline for data center environmental monitoring. The pipeline comprised a configurable sensor and fog layer, a scalable cloud backend split across two independently deployed Lambda functions, and deployment and testing on a public cloud platform. The project also aimed to critically analyse the architectural decisions taken at each layer against the alternatives available. All three layers were implemented, tested, and deployed to a real AWS account rather than a local simulation alone.

The most significant finding of this project is not architectural but methodological: local testing against a cloud emulator, however thorough, is not a substitute for testing against the real target platform. Both defects discovered during the real AWS deployment, the credential-handling bug and the DynamoDB pagination bug, passed the full local test suite and a full LocalStack-backed integration test cleanly, and would have shipped undetected had the project stopped at the local-simulation stage. The live load test's eventual-consistency finding reinforces the same point from a different angle: even a passing local test cannot exercise a platform behaviour, like SQS's approximate-count consistency model, that the local emulator does not itself reproduce.

The most difficult part of this project was not writing application code but discovering the AWS Academy Learner Lab account's platform and IAM constraints. Several of these, including the shared `LabRole` and an account lifetime that was session-based, not persistent, shaped design decisions outright; they were not incidental factors. Future work would include converting the fog and sensor tier itself to a fully event-driven, scheduled-Lambda model in place of a continuously running EC2 instance. If the account restrictions were lifted, future work would also provision two separate least-privilege IAM roles for `dce-processor` and `dce-api` instead of the single shared `LabRole` this deployment was constrained to.

#### REFERENCES

- [1] C.-J. M. Liang, J. Liu, L. Luo, A. Terzis, and F. Zhao, "RACNet: A high-fidelity data center sensing network," in Proc. 7th ACM Conf. Embedded Networked Sensor Systems (SenSys '09), Berkeley, CA, USA, 2009, pp. 15–28.
- [2] M. G. Rodriguez, L. Ortiz Uriarte, Y. Jia, K. Yoshii, R. Ross, and P. H. Beckman, "Wireless sensor network for data-center environmental monitoring," in Proc. 2011 Fifth Int. Conf. Sensing Technology (ICST), Palmerston North, New Zealand, 2011, pp. 533–537.
- [3] P. Mell and T. Grance, "The NIST Definition of Cloud Computing," National Institute of Standards and Technology, Gaithersburg, MD, USA, Special Publication 800-145, 2011.
- [4] F. Bonomi, R. Milito, J. Zhu, and S. Addepalli, "Fog computing and its role in the Internet of Things," in Proc. 1st ACM Mobile Cloud Computing Workshop (MCC '12), Helsinki, Finland, 2012, pp. 13–16.
- [5] J. Giacomoni, T. Moseley, and M. Vachharajani, "FastForward for efficient pipeline parallelism: A cache-optimized concurrent lock-free queue," in Proc. 13th ACM SIGPLAN Symp. Principles and Practice of Parallel Programming (PPoPP '08), 2008, pp. 43–52.
- [6] L. Lamport, "Concurrent reading and writing," *Commun. ACM*, vol. 20, no. 11, pp. 806–811, 1977.
- [7] P. T. Eugster, P. A. Felber, R. Guerraoui, and A.-M. Kermarrec, "The many faces of publish/subscribe," *ACM Comput. Surv.*, vol. 35, no. 2, pp. 114–131, Jun. 2003.
- [8] A. Carzaniga, D. S. Rosenblum, and A. L. Wolf, "Design and evaluation of a wide-area event notification service," *ACM Trans. Comput. Syst.*, vol. 19, no. 3, pp. 332–383, Aug. 2001.
- [9] G. DeCandia et al., "Dynamo: Amazon's highly available key-value store," in Proc. 21st ACM SIGOPS Symp. Operating Systems Principles (SOSP '07), 2007, pp. 205–220.
- [10] M. Shahrad, J. Balkind, and D. Wentzlaff, "Architectural implications of function-as-a-service computing," in Proc. 52nd Annu. IEEE/ACM Int. Symp. Microarchitecture (MICRO-52), Columbus, OH, USA, 2019, pp. 1063–1075.
- [11] A. Akbulut and H. G. Perros, "Software versioning with microservices through the API gateway design pattern," in Proc. 2019 9th Int. Conf. Advanced Computer Information Technologies (ACIT), Ceske Budejovice, Czech Republic, 2019, pp. 289–292.